

Advanced Concepts of Machine Learning: Sparse Autoencoder

Kevin Trebing (i6192338), Alberto Perez (i6201157)

November 5, 2018

1 How to use

Execute the python file *sparse_autoencoder.py* with python3. An optimization will run that trains the autoencoder. Different parameters can be chosen, such as the learning rate α (called ALPHA) with default value 0.1 and λ (called LAMBDA) with default value 0.001. Furthermore, the training stops when the total error of one epoch (i.e. seeing each training sample once) is below a certain threshold (default for this is 0.0001).

We have some issues in our code that cause the hidden layers activation not to converge to proper zeros and ones. A cause for this is probably the error calculation of the output layer. We use the summed up error to tell whether we are converged or not. Since the error values there can be positive and negative, a sum is not appropriate to calculate the error. But using the absolute values leads to no convergence at all, which means that we have an error in our back-propagation. Since we already spend about 12 hours on this we left it at the "wrong" convergence definition.

2 Learning performance

The network needs about 15.000 – 20.000 epochs to converge. Note that it still does not learn the mapping perfectly, instead of zeros it most often leads to values around zero. The same goes for values of one, here it leads to values around one. The time until convergence is about 10 seconds.

3 Parameters

For higher values of alpha, it is more likely that the system will converge at the wrong stability point. For high values, decreasing the size of the error used for determining convergence can decrease the chances of wrong stability points. However this also requires many iterations to achieve convergence.

The highest alpha value which admits for a reasonable amount of convergence is 5. For such a value, convergence does not take more than 2000 epochs, however sometimes the loop can go on for indefinite time. Reducing the value of alpha to around 0.8 gives a more robust probability of convergence. However obtaining this convergence usually takes longer. To reduce the time taken a regularisation term is added; this is represented by the Lambda factor.




Figure 1: Different alpha values lead to different convergence speeds. The regularization factor was fixed at 0.001.

Lambda is not very useful for large values of alpha (where convergence is quick already). However for small values of alpha (<0.8), lambda can lead to a faster convergence (for example reducing the epochs necessary by one order of magnitude). This does incur a penalty as the network loses some preciseness. However, as long as the lambda is not too large (<0.001), the output values of the network correctly identify the input values.

4 Interpretation of the learned weights

The weight matrix going from the input layer to the hidden layer has 27 weights (9×3) and the weight matrix going from the hidden layer to the output layer has 32 weights (4×8). This makes a total of 59 weights. From inspection approximately half of the value in the weights will be negative (but not necessarily exactly half of them are negative).

In a forward pass from one layer to the next, the magnitude of each weight indicates how much importance to give to each value from the input layer (of that pass). The sign indicates whether the input is adding positively or subtracting from the output value. Each weight is used once in one forward pass from one layer to the next, as there is one weight for each possible combination of inputs and outputs. In the backward propagation the magnitude of the weights indicate how much importance to give to the error found at the output activation values

(relative to each original input node).

In the forward pass all the input activations (and the bias) are multiplied by the weights going to a specific output activation (and incoming from each different input), i.e. for activation a_3 , it would be weights $W_{3i} \dots W_{3n}$. However in the backward pass the errors are multiplied by the weights indicating they originate from a specific node (while the errors are incoming from each different output node), i.e. for the error at activation a_2 , it would be weights $W_{i2} \dots W_{n2}$.

The weights are also used when calculating the regularisation. The effect of the regularisation is to try and promote the current behaviour of that weight. A large positive weight will have its value increased, while a large negative weight will have its value decreased. A similar effect will be seen in smaller weights but at a smaller scale. The overall effect is a reduction on the changes the update might inflict opposing the current pattern of behaviour of the weights. The regularisation tends to create weights with a slightly smaller magnitude when compared to same run with no regularisation.

When looking at the three neurons in the hidden layer we can see that an (almost) binary representation was learned. A human interpretation could be that if a neuron has an activation bigger than 0.1 we regard it as firing. Then we have a perfect binary encoding of the input.